

# What is Cassandra?

The Apache Cassandra database is the right choice when you need scalability and high availability without compromising performance. Linear scalability (<http://techblog.netflix.com/2011/11/benchmarking-cassandra-scalability-on.html>) and proven fault-tolerance on commodity hardware or cloud infrastructure make it the perfect platform for mission-critical data. Cassandra's support for replicating across multiple datacenters is best-in-class, providing lower latency for your users and the peace of mind of knowing that you can survive regional outages.

## PROVEN

Cassandra is in use at Constant Contact (<http://www.slideshare.net/daveconnors/cassandra-puppet-scaling-data-at-15-per-month>), CERN (<http://planetcassandra.org/blog/post/cassandra-at-cern-large-hadron-collider/>), Comcast (<http://www.slideshare.net/planetcassandra/nyc-tech-day-using-cassandra-for-dvr-scheduling-at-comcast>), eBay (<http://www.slideshare.net/jaykumarpatel/cassandra-at-ebay-13920376>), GitHub (<http://planetcassandra.org/blog/post/analytics-at-github-with-apache-cassandra/>), GoDaddy (<http://planetcassandra.org/blog/post/godaddy-worlds-largest-domain-name-registrar-and-web-host-provider-utilizes-cassandra-for-replication-and-scalability/>), Hulu (<http://planetcassandra.org/blog/post/cassandra-used-to-build-scalable-and-highly-available-systems-at-hulu-streaming-content-to-over-5-million-subscribers/>), Instagram (<http://planetcassandra.org/blog/post/instagram-making-the-switch-to-cassandra-from-redis-75-instasavings/>), Intuit (<http://www.slideshare.net/planetcassandra/3-mohit-anchlia>), Netflix (<http://www.slideshare.net/adrianco/migrating-netflix-from-oracle-to-global-cassandra>), Reddit (<http://planetcassandra.org/blog/post/reddit-upvotes-apache-cassandras-horizontal-scaling-managing-17000000-votes-daily/>), The Weather Channel (<http://planetcassandra.org/blog/post/make-it-rain-apache-cassandra-at-the-weather-channel-for-severe-weather-alerts/>), and over 1500 more companies (<http://planetcassandra.org/companies/>) that have large, active data sets.

## FAULT TOLERANT

Data is automatically replicated to multiple nodes for fault-tolerance. Replication across multiple data centers is supported. Failed nodes can be replaced with no downtime.

## PERFORMANT

Cassandra consistently ([http://vldb.org/pvldb/vol5/p1724\\_tilmanrabi\\_vldb2012.pdf](http://vldb.org/pvldb/vol5/p1724_tilmanrabi_vldb2012.pdf)) outperforms (<http://www.datastax.com/resources/whitepapers/benchmarking-top-nosql-databases>) popular NoSQL alternatives in benchmarks and real applications (<http://blog.markedup.com/2013/02/cassandra-hive-and-hadoop-how-we-picked-our-analytics-stack/>), primarily because of fundamental architectural choices (<http://www.datastax.com/dev/blog/2012-in-review-performance>).

## DECENTRALIZED

There are no single points of failure. There are no network bottlenecks. Every node in the cluster is identical.

## SCALABLE

Some of the largest production deployments include Apple's, with over 75,000 nodes storing over 10 PB of data, Netflix (2,500 nodes, 420 TB, over 1 trillion requests per day), Chinese search engine Easou (270 nodes, 300 TB, over 800 million requests per day), and eBay (over 100 nodes, 250 TB).

## DURABLE

Cassandra is suitable for applications that can't afford to lose data (<http://wiki.apache.org/cassandra/Durability>), even when an entire data center goes down.

## YOU'RE IN CONTROL

Choose between synchronous or asynchronous replication for each update. Highly available asynchronous operations are optimized with features like Hinted Handoff (<http://wiki.apache.org/cassandra/HintedHandoff>) and Read Repair (<http://wiki.apache.org/cassandra/ReadRepair>).

## ELASTIC

Read and write throughput both increase linearly as new machines are added, with no downtime or interruption to applications.

## PROFESSIONALLY SUPPORTED

Cassandra support contracts and services are available from third parties (<http://wiki.apache.org/cassandra/ThirdPartySupport>).

© 2016 The Apache Software Foundation (<http://apache.org>). Apache, the Apache feather logo, and Apache Cassandra are trademarks of The Apache Software Foundation.